

Linked VISTA

A Metadata-Centric Approach to VA Master Data Management

February 24, 2015
Pacific Joint Information Technology Center
Maui, HI

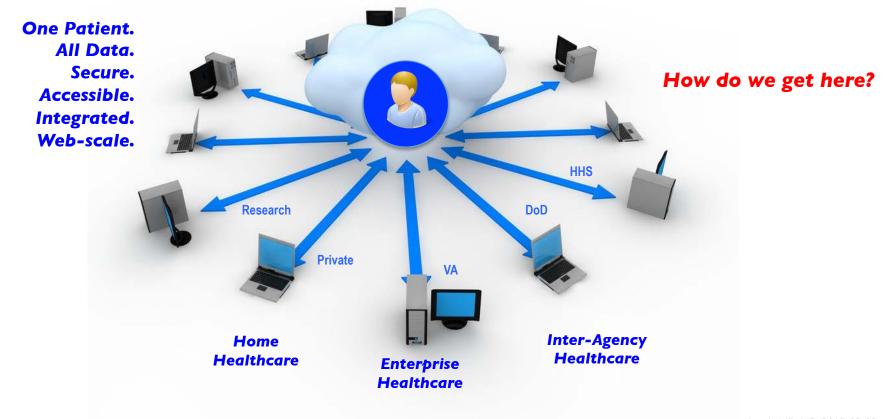
Rafael Richards MD MS

VHA Office of Informatics and Analytics U.S. Department of Veterans Affairs



Vision: Linked VISTA

A Linked Data Approach to Patient-Centric Care





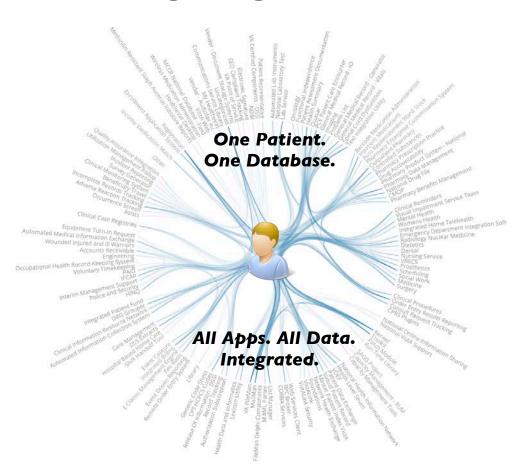
VISTA Overview

- VISTA Architecture: Single-integrated EHR
- VA Enterprise: One Patient. Many VISTAs
- Problem: Isolated VISTAs. Fragmented data.
- Approach: Linked VISTAs.
- Vision: OneVA.



VISTA Architecture: Single-integrated EHR

A Patient-Centric Single integrated EHR



Patient-centric Health Record

The data architecture of VISTA consists of over 160 modules for clinical care integrated within a single common multidimensional data engine (MDE).

In VISTA, both business logic (Applications) and data (Database) are managed within the multidimensional data engine. This provides the tight integration of applications to data, and to a single common integrated database.

The integration between VISTA applications (outer ring) and shared data (inner circle) is visualized, showing the shared data flow between applications.



VA Enterprise: Many VISTAs

1200 VA healthcare facilities throughout the US Supported by 130 VISTA systems



Care coordination for patients within VA requires integration of patient data across all 130 VISTA systems.





Problem Statement

VA is comprised of 130 unique VISTAs. This fragments patient data and care.



Problem: Isolated VISTA

VISTA Silos



Many Databases.

















Many VISTAs. Many Models. Fragmented Data.

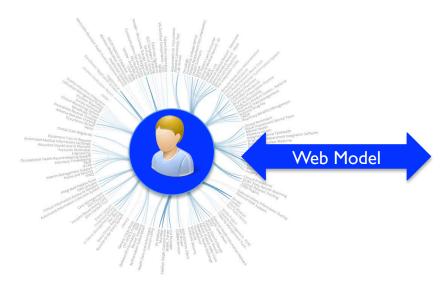
While each VAVISTA system is individually highly integrated, each system has its own distinct data model (shown as different color).

Different data models creates data silos – fragmenting patient data and fragmenting patient care.



Approach: Linked VISTA

Exposing VistA's Data Model securely as a Webstandard model enables web-scale integration.



Web-standard
Integration requires:

Common Data Model.
Data Security.
Data Services.

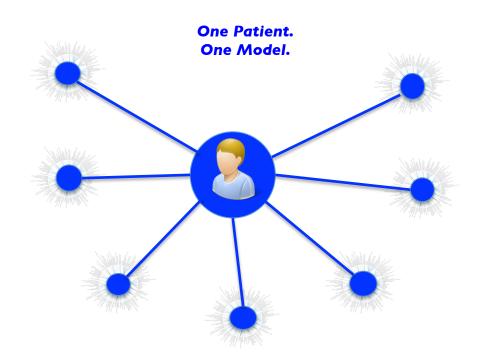
A model-centric approach to VISTA data modernization would provide comprehensive granular, metadata driven access and security to all data.

This would allow publishing, linking, and exchanging, with patient-centric security, granular data with internal and external systems with international Web standards.



Approach: Linked VISTA

Linked VISTA



Many VISTAs. One Model. Integrated Data.

Exposure and cross-linkage of VISTA's local models to a common VISTA model provides an integrated view of all patient data from all VISTA systems.





VISTA Data Model (VDM)

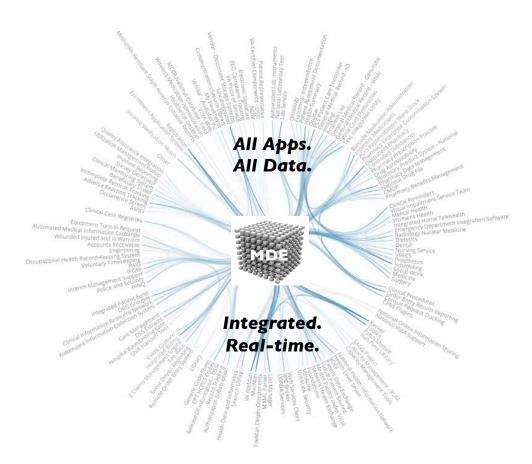
A Path to VISTA Data Management:

- VISTA's Database
- VISTA's Data Model
- VISTA's Data Model Exposed
- Benefits of Leveraging VISTA's Data Model:
 - Master Metadata Management
 - Centralized Knowledge Management
 - Master Data Definition
 - Patient-Centric Security Model
 - Separating Business logic from Data
 - Query Access



Review: VISTA's Database

The foundation of VISTA is a high performance Multidimensional Dataintegrated Application Engine in which all data and all applications are fully integrated in real-time with each other and to one single authoritative data source.



VISTA's integrated application data engine. All 160+ applications are integrated with their data and logic inside the multidimensional data engine (MDE). This keeps transactional patient data and logic highly integrated for real-time use within one single data store.

VA uses the same healthcare industrystandard data engine that nearly 50% of hospitals in the US currently use as their core EHR database.

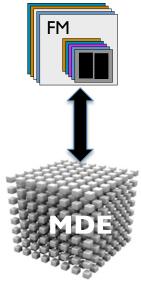
This same multidimensional data engine is also used as the real-time transaction engine of the world's five largest banks, the largest Wall Street trading systems, and over half a million ATMs in North America.



Review: VISTA's Data Model (VDM)

All real-time transactional operations in VISTA take place within the the multidimensional data engine (MDE). To provide consistent structure and model for the data, a data dictionary driven hierarchical data storage system is overlaid on top of the MDE. All VISTA applications read and write data to this hierarchical store using a file manager called Fileman.



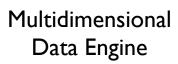


Hierarchical Graph Store



This layer provides the data structures, data models, and data management for VISTA.
Fileman is the DBMS for this hierarchical store.

Apps Data





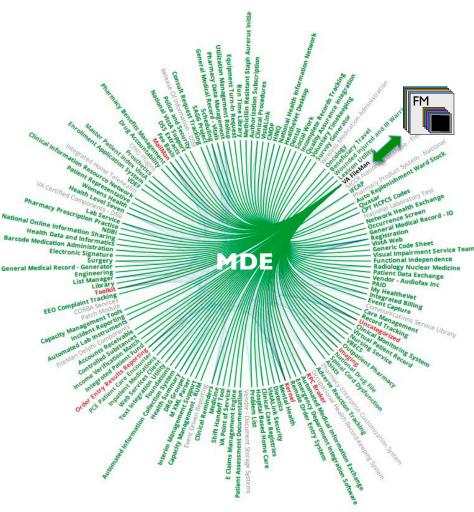
This is the data-structure flexible transactional core of VISTA. This core is a healthcare industry-standard for nearly 50% of the hospitals in the US. This is also used as the real-time transactions engine of the world's five largest banks, the largest Wall Street trading systems, and over half a million ATMs in North America.





Review: VISTA's Data Model (VDM)

All VISTA applications read and write data to VISTA's data engine through the NoSQL hierarchical data manager, Fileman (FM). Unlike many NoSQL databases, VISTA's database is not schema-less, but schema-driven. VISTA's schema is self-documenting through data dictionaries in Fileman. Surprisingly, this fact has not yet been leveraged to expose and modernize the data model.



All VISTA data is organized through its data model, managed by Fileman.

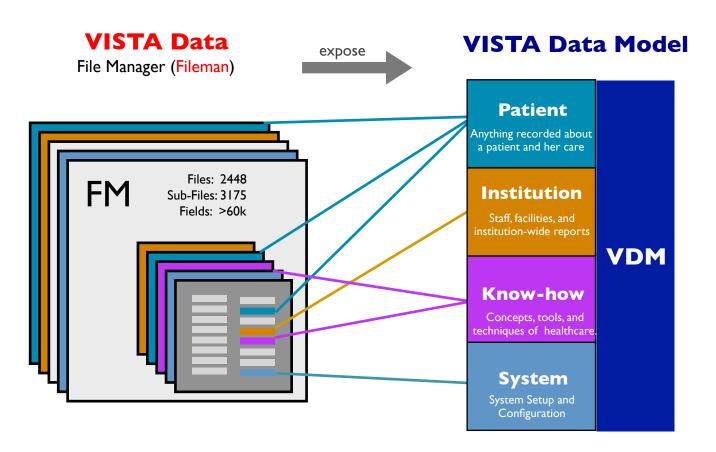
All VistA applications use Fileman to access, structure, and query all data in VistA.

The green lines represent the data flow of all VISTA application to Fileman for all read and write operations to the Multidimensional Data Engine (MDE).



VDM: Expose and Leverage the Model

The first stage of data modernization is to expose and leverage VISTA's real, live operational data model. Since this is just metadata, there is no patient data involved. Unlike many NoSQL databases which are schema-less, VISTA's NoSQL database model is self-documenting through Fileman. This allows us to render this in a standard definition format. In this new webstandard medium, data can be sorted, tagged, searched, and organized by data categories such as by patient, institution, know-how, or system information.



Access: Thousands of RPCs, API's, HL7

Model: Unknown

Access: Single query access

Model: Consistent, Transparent

RM Richards MD MS 2015-02-23



VDM: Master Data Management

A benefit of a VISTA data model allows one to manage data logically across all VISTA application boundaries independent of the source of the data. This lets one manage data with much flexibility, including logically partitioning and managing the data using metadata tagged categories (such as Patient, Institution, Know-how, and Systems information).

VISTA Data Model

- One may logically partition data by any class of data, such as Patient data, Institution data, reusable Know-how, and System configuration data.
- One may extract and move <u>all</u>
 patient data from system to
 system with one operation,
 making system configuration
 migration and patient record
 movement far more efficient.
- One may apply security metadata or protocols to any of these logical classes of data.
 For example a patient-centric security model for patient data, and thus enforce patient-centric controls on information exchanged.

Patient Anything recorded about a patient and her care Institution Staff, facilities, and institution-wide reports Know-how Concepts, tools, and techniques of healthcare. System System Setup and Configuration

Benefits:

Patient Data Management

Extract and manage patient data with patient-specific security and metadata, allowing patient-centric controls on data access and exchange.

Institutional Data Management:

Institution specific data can be exchangeable and centrally manageable

Knowledge Management:

Common medical concepts, standards, and know-how may be identified and managed as a clearly defined class of VISTA data.

System Management:

The entire configuration of a system can be identified, extracted, and transported, and inserted from system to system



VDM: Patient-Centric Data Security

An exposed VISTA data model allows one to tag and partition certain classes of data separately from all other data in VistA. Specifically, this allows one to granularly partition any and all Patient data from all other kinds of data in VistA. This provides true, direct, "on the metal" security on patient data itself.

VISTA Data Model

- The most important class of data to apply security is the Patient data category. This will allow very granular patient-centric security on the data itself.
- Current VISTA security is based on actions one can take using a legacy menu system - which has no relationship to the data.
- With a data-centric security model, we can specify not just what type of data ("Mental Health Record"...) but whose data ("For patient X").
- This is much more specific and secure than the prior VistA security model.



Patient-centric security model

Extract and manage patient data with patientspecific security and metadata, allowing patient-centric controls on data access and exchange.

This is compatible with security notions in Meaningful use Data Exchange: it can suppress even data that exists if there is no access permission.

It is essential to improve precision in data security to permit access to VISTA data securely. Otherwise one will have to reverse-engineer 3300 legacy RPCs and their one-off use of Kernel's menu options for each payload.

A patient-centric security model is much more appropriate, flexible, and secure as a foundation for patient data security than the current VistA security model. The current VistA security model provides security only <u>indirectly</u>, through legacy controls of a <u>menu system for a legacy roll and scroll terminal interface — which has nothing to do with the type of data at all (!).</u>



VDM: Analytics Driving Interoperability

Comprehensive exposure and analytics of the VISTA Data Model will drive enhanced data use and interoperability as well as a major improvement in the structure of the database itself. To address these and other areas, focused reports could be generated from the model including:

Report	Activity	
Inconsistencies between VistA data models	Drive dictionary and code fixes in various centers so that every center is running the same consistent model	Enterprise Data Model
Isolate centrally and locally managed "know-how"	Enables the next generation of enterprise knowledge services that seamlessly synchronize VistAs and other applications	Centralized Knowledge Management
Under-definition in the model	Too many ill-defined string values and not enough nuance ("zip code", "telephone number") can be defined, and provided additional metadata ("home", "work", "mobile") leading to a plan for incremental dictionary improvement	Enterprise Data Definition
Key logic performed within FileMan	The barrier between the data store and business logic will be laid bare. This report will encourage the movement of certain types of logic into FileMan and out of less maintainable procedural code.	Clean separation of business logic from data
Overlooked but highly valuable types of patient data	Improve VistA Data mining (for CDW etc.) and interoperability (more comprehensive electronic patient records).	Clinical Research
(Some of this may have been hard-set by the application logic, and overlooked by Fileman or DD)	Without a complete model, how do you know what you're leaving behind?	Interoperability
VistA model/ FHIR comparison (key types)	Show how a direct from VistA transformation can remove the need for redundant intermediate, hard to maintain procedural code	Accelerate Data Exchange
Isolate patient from other types of data	Enable patient-data access control rather than the crude option/API security now in VistA	Patient-centric Security Model



Linked VISTA

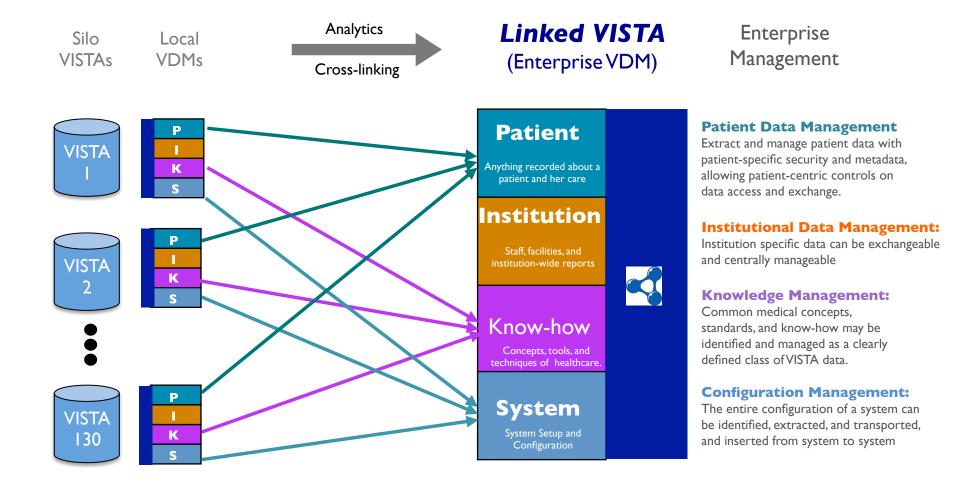
Features

- Enterprise Cross-linked VISTA Model
- Comprehensive Data Access
- Web-standard representation language
- Web-scale semantic integration
- Knowledge Discovery
- Native terminology integration



Linked VISTA: The Enterprise Cross-Linked VISTA Data Model

Exposing VistA's Data Model (VDM) and applying analytics allows one to cross-link all local VDMs to an enterprise VISTA data model (Linked VDM), providing the capability for Enterprise query and Enterprise data management. This leverages the capability of Linked Data to create a cross-VISTA (enterprise) data model.



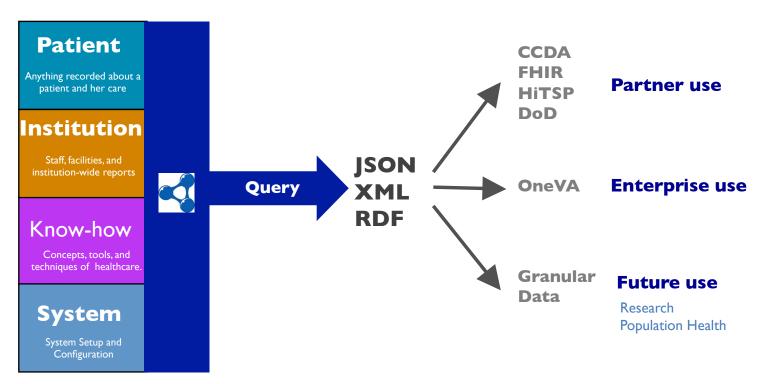


Linked VISTA: Comprehensive Access

One benefit of creating a cross-VISTA enterprise data model (Linked VISTA) is that it can be queried against <u>any VISTA</u> for <u>any data</u> with with <u>one</u> web-standard query.

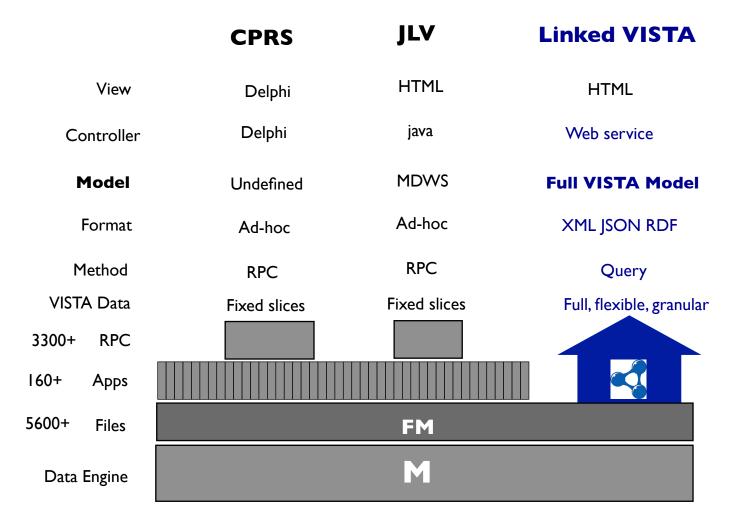
This would allow any authorized system to securely and directly query authoritative VISTA data in real-time with one standard query language. The output of these queries will be in all modern web-standard forms to maximize secondary use.

Linked VISTA





Linked VISTA: Comprehensive Access



VISTA Data Model (VDM) can access all data spanning 160+ applications with full granularity because the fully exposed VISTA Data model bridges all applications. No legacy APIs, HL7, RPCs, or MUMPS. Just VISTA data. All of it. Real-time.

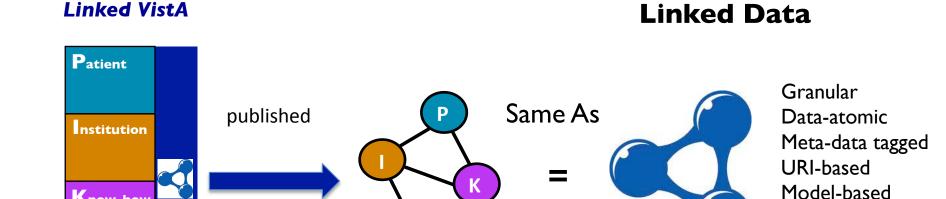


Know-how

System

Linked VISTA: Web-standard Representation

Use of web-standard representation maximizes data re-use, meshing, and re-mixing with the maximum number of other sources of data for research and patient care.



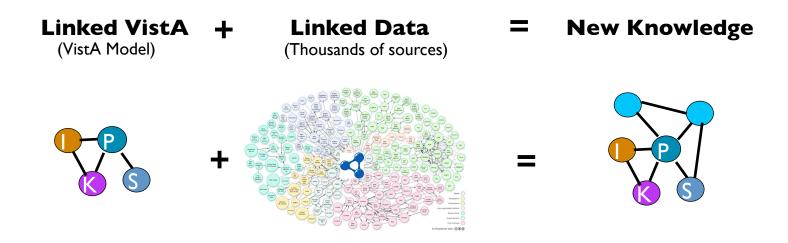
Linkable

Web Standard



Linked VISTA: New Knowledge Discovery

Linked VISTA can participate in federated queries over unlimited number of other Linked Data sources, enabling meshing, enrichment, and ultimately, new knowledge discovery.

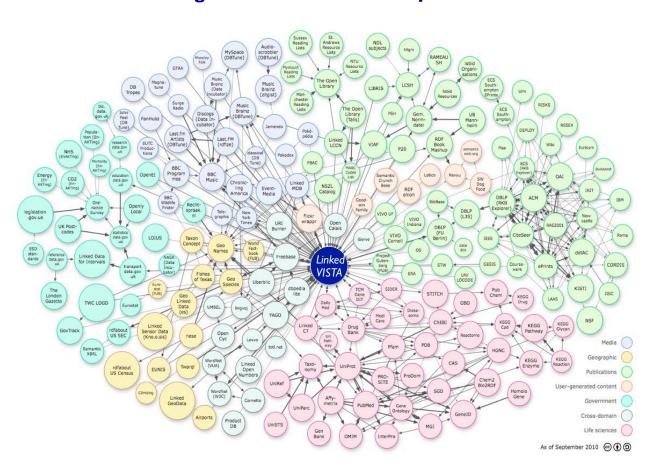


Because data in the VistA Data Model can be represented just like any other Linked Data resources, one can mesh VistA data directly with unlimited sources of internal or external, public or private life sciences, and other scientific or healthcare related data sources. This leads to discovery of new relationships between different sources of data - and new knowledge.



Linked VISTA: Internet-scale semantic integration

Representing VistA data in a Linked Data form supports real-time semantic integration with thousands of other other linked sources.



Linked VistA

VistA Data - represented in standard Linked Data form - can be interlinked with any and all other Linked Data sources.

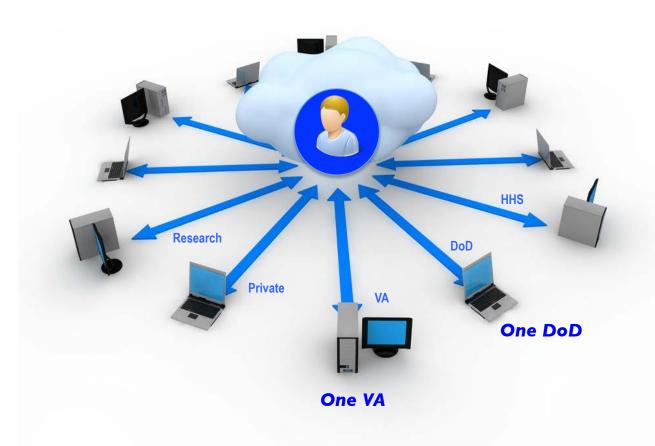
This will enable meshing, enrichment, and augmentation of patient data with all other sources, providing a comprehensive view of all patient data from all locations, clinics, hospitals, or the home.

Data sources also include patientgenerated, mobile device, TeleHealth, and any other Internetenabled device data (Internet of Things).



Vision: Linked VISTA

Web-standards approach to provide Patient-Centric Collaborative Care



One Patient.
One Model.
Integrated.
Web-scale.



Reference

Linked Data Summary Linked Data Links

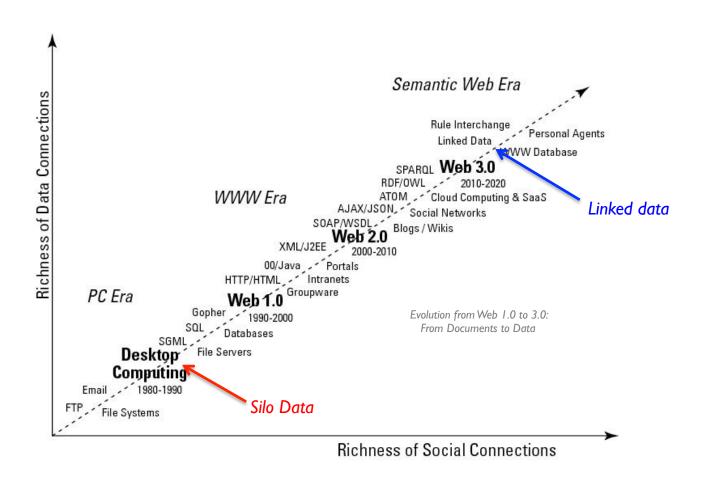


Linked Data

- Evolution of Computing: Increased connectivity
- What is it Linked Data?
- What problems does it solve?
- Who uses Linked Data?
- Health Data: Many diverse models
- Linked Data: Accommodates model diversity
- Health Data: PCAST Recommendation



Evolution of Data management: Increased Connectivity

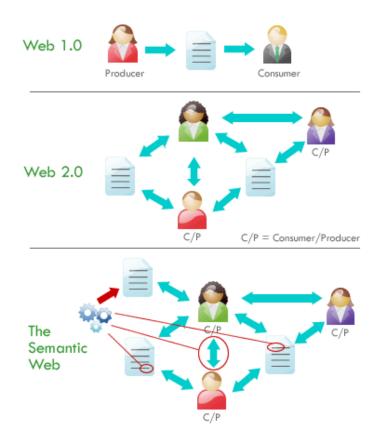


DHCP was officially launched in 1982 and is the foundation of all VA VISTA and DoD CHCS systems.

This was in the early PC Era, before the age of inter-networked data ("the internet"), when databases were isolated inside institutional systems (Silo Data: red arrow). Note that the TCP/IP internet protocol had not even been invented at that time. To bring any VISTA system into today's WWW Era, allowing data to be linked to all VISTA systems within VA enterprise (or to any CHCS system in DoD), VISTA needs be web data-standards capable. This would transform VISTA to a WWW Database that could manage Linked Data (blue arrow).



Evolution of the Web: From Documents to Data



Web 1.0: Document Web (HTML)
Linked Documents

Read-only web (humans only)

Web 2.0: Social Web

Linked People

Read-write web (humans only)

Web 3.0: Semantic Web (RDF)

Linked Data

Read-write web (machine processable)



Linked Data: What is it?

The World Wide Web (W3C) Standard for semantic information integration



HTML (hypertext markup language)
For **humans** to exchange information



Linked Documents (Document Web)



RDF (resource description framework)
For **computers** to exchange information



Linked Data (Semantic Web)



"The Semantic Web [Linked Data] provides a common framework that allows data to be shared and reused across application, enterprise, and community boundaries."

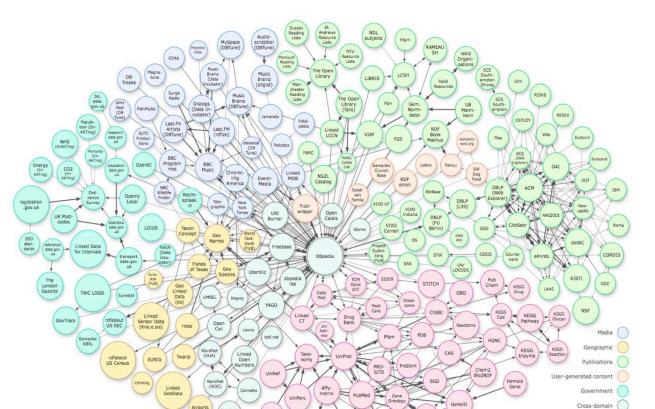
Tim Berners-Lee, MIT Professor and Inventor of the World Wide Web

As a W3C standard this supports Internet-scale data integration.



Linked Data: What does it enable?

Web-scale semantic integration of data





Linked Data

This figure shows the Linked Open Data (LOD) cloud, which semantically links hundreds of Linked Data sources including Media, Geographic, Government, and Life Sciences databases.

Each circle represents one data source or database. These are semantically linked to other data sources, creating a single virtual federated internet-scale database.

At the center of is DBPedia, the Linked Data version of Wikipedia, which is semantically linked to hundreds of data sources.

Life sciences

As of September 2010 @ 10

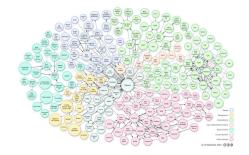


Linked Data: Who Uses It?

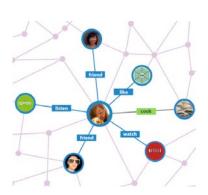
The Linked Data approach to **Internet-scale semantic data integration** is used by the worlds largest organizations such as Google, LinkedIn, Facebook, and IBM Watson.



Wikipedia (knowledge graph)



IBM Watson (knowledge graph)



LinkedIn (professional graph)



Facebook (social graph)







Linked Data: Accommodates both Standards and Innovation

Information Models:

An apparent conflict between standardization and innovation

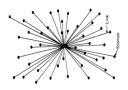
Standards: need to remain static in order not to be disruptive for adopters.

- Static, Brittle
- Centralized
- General (Common Denominator)
- Committee-driven
- Large, "all-or-nothing", disruptive updates

Innovation: requires continuous evolution of thousands of new information models.

- Adaptive / Evolutionary
- Decentralized
- · Highly Specialized, "Best of breed"
- End-user / specialist driven
- · Small, continuous, low-impact updates

What are the options?



Centralized, Model-rigid approach: For exchange of information to occur all models must remain fixed, and data must go through only one central 'broker' model..

Our current approach to healthcare data

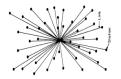


Decentralized, Model-flexible approach: Multiple models peacefully co-exist and evolve, mediated by their ability to freely link to any model at all times. In this approach, all models are free to evolve AND are capable of resolving to a common standard model at all times. The only current technology that supports this granular data-oriented approach is RDF (Linked Data).

Linked Data supports <u>both</u> standardization and innovation



Data Integration: Legacy vs. Linked Approach





Architectural Issues

Current Approach

Linked Data (RDF)

Data model characteristics	Model-rigid. Only one lowest common denominator model unifies information. Must remain unchanged to orchestrate. Restrictive expression.	Model-flexible. All data models may independently evolve. Maximizes expressivity.
Data model compatibility	No model diversity permitted A one-size-fits-all mega-model	Multiple models peacefully coexist Data model flexible
Data model evolution	Costly and difficult to evolve models. Due to model-rigid architecture.	Cheaper and easier to evolve models. Due to model-flexible architecture.
Data access method	Downloading + Aggregating	Linking + Federating
Scalability: incremental effort required to add new data sources	Common model must be updated	Individual models may be independently and incrementally semantically linked.
Primary Function	Data Syntax and Transport	Data Model
Granularity	Document-centric (message-centric)	Data-centric (individual data elements)
Semantics	Weak semantics. Extrinsic, separate from the data. Depends on an external data model.	Strong semantics. Intrinsic, integrated with the data.



Links: Linked Data

W3C Linked Data Standard

http://www.w3.org/standards/semanticweb/data

W3C Linked Data Platform

http://www.w3.org/TR/ldp/

W3C Semantic Web Healthcare and Life sciences

http://www.w3.org/blog/hcls/

HL7-RDF Healthcare Standards Work Group

http://wiki.hl7.org/index.php?title=RDF_for_Semantic_Interoperability

Semantic Web

http://en.wikipedia.org/wiki/Semantic Web

Linked Data: Tools (>1000)

http://www.mkbergman.com/sweet-tools

Linked Data: Adaptors (100's)

http://www.w3.org/wiki/ConverterToRdf

Linked Data: Roadmap for Healthcare Interoperability

http://yosemiteproject.org



VHA: RDF is an approved technical reference model

